

Golang Compiler

Final Report

Group 2

April 15, 2016

Members: Shabbir Hussain
Ossama Ahmed
Michael Ho

1 Introduction

Our semester project was to build a go programming language compiler. Our compiler takes in source code from a subset of go called golite and compiles source code to jasmin code. The project requirements were to create a multi-stage compiler which:

- Scans source code for tokens and keywords
- Parses the tokens according the grammar of golite into an Abstract Syntax Tree (AST)
- Weeds the AST to comply with the grammar of golite
- Type checks the source according to the typechecking rules of golite
- Emits a pretty printed version of the source code
- Emits a symbol table of the type-checked source
- Code generates code in the target language (jasmin).

All the specifications of the golite language can be found at <http://www.sable.mcgill.ca/hendren/520/2016/>. This report will explain how we achieved the above mentioned requirements.

2 Compiler Structure

Our compiler is built using the OCaml Lex and Menhir toolkit. We chose Ocaml for its strongly typed language in hopes to catch errors early in the development process. Furthermore, Ocaml contains many online help resources. Lastly, it is a great opportunity for hone our knowledge and programming skills in Ocaml. Ocamllex is tool for scanning the source into tokens which are encoded as Ocaml types. Menhir is the parser tool used to verify the syntax of the source code. The rest of the compiler is built using Ocaml native libraries. A high level overview of all the Ocaml modules, their inputs and their dependencies can be seen in figure 1. One of the design difficulties was the tight coupling of the AST to all the modules. This meant that if any changes were needed in the AST then all modules interacting with the AST had to change. As the project progressed, we often went back and changed the AST to add more information. The rest of this section dives deeper into the design decisions of each module.

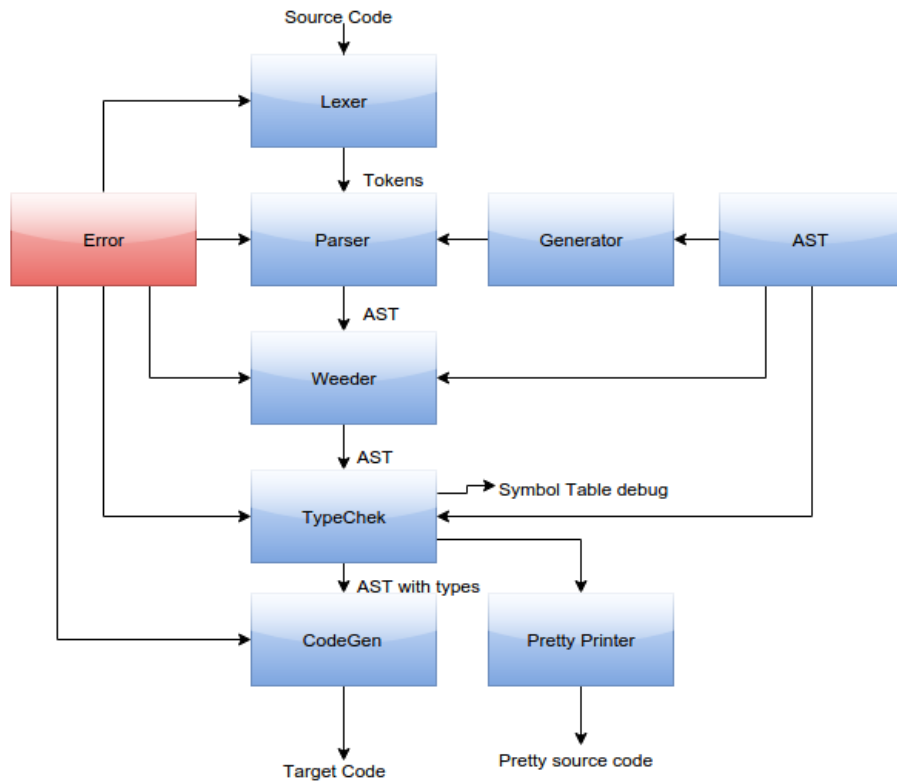


Figure 1: Overview of the modules in the compiler design

2.1 Error Module

We decided to have error messages standardized using an error module. By doing this all our error messages are of type `GoliteError` and contain the line number at which the error occurred.

2.2 Abstract Syntax Tree

The abstract syntax tree is composed of Ocaml Data Type. It is very close to the original golite grammar. We quickly realised that it would be advantageous to have each node hold the line number at which it occurred. This allowed us to give more detailed error messages when a parsing error occurs. Subsequently, it was useful in all the other phases as well or error checking. During type checking, we also added a field to every expression to contain type information.

2.3 Parsing Phase

In order to have less dependency on the AST we decoupled the AST node constructors using a generate module. The generate module provided a standard interface for constructing the AST. The generator module proved its usefulness when we added a type field to the constructors of the AST. We only needed to change the generator code and not the parser. We initially construct the AST having all the type information set to empty in the generator module. One disadvantage of the generator is that it adds complexity to our code. To save time we considered the source code of our compiler as our documentation. However, to understand the AST received in the type-checker, we need to cross reference the AST with the parser and the generator.

2.4 Weeder

We decided to place some syntactic checks in the weeding phase as opposed to the parser because it is easier to check. For example it is easier to check if a switch statement has only one default case after all the cases have been defined. The parser builds the AST bottom up and our weeder checks the syntax from the top down. It searches through every node until it hits a node that needs to propagate information to its leaf node. For example, a loop node needs to pass on information to its children that it is in a loop such that we can catch dangling break and continue statements.

2.5 Type checking Phase

The type checker uses the same tree searching algorithm as the Weeder. We've implemented the type checking based on the specification on the sable course page as well as the reference compiler. We initially implemented the type checking as a two pass compiler in order to have function declarations with global scope since the go language allows it. We later removed this feature when verifying with the reference compiler. We added types to expressions in the AST

such that we can print them for debugging purposes as well as code generating appropriate code. We decided to create constructors for each type to be included in the symbol table; this allows us to have a common interface for the symbol table. We also created a recursive type called `symType`. This allows us to create a "type" for structs which have structs in them.

2.6 Pretty Printer

For Pretty Printing, we decided to rewrite the `PrettyPrinter` using OCaml's imperative style (begin / end statements) and printing to `stdout`. Instead of printing as one large string as we did in milestone 1, now with imperative printing, it became easy to implement `-pptype` and soon code generation should be easier as well.

for the `-dumpsymtab` flag, the `symtab` file will be in the same folder as the executable for the `-pptype` flag, the `pptype` file will be in the same folder as the source file

2.7 Code Generation

Our target language is `jasmin`. We initially chose `jasmin` because we believed we may be able to use the `peephole` optimizer in our compiler. Although it would have been interesting, this was not possible because we wanted to emit code that uses the full `jasmin` language; the `peephole` optimizer does not support floating point numbers. None the less, the exercise of the `peephole` optimizer allowed us to familiarize ourselves with `jasmin` code.

We added some helpers to enable us to generate code to `jasmin`. We had a functions to

- to start and end scope
- to reference a variable name to a local
- to get type checking information from the AST

Futhermore, some of the functionality of `golite` does not translate directly into `jasmin`. We had to implement some macros and algorithms in `jasmin` to accommodate these functionalities. The rest of this section will show detailed examples of these algorithms.

Structs

We treat structs as java classes. Any time we encounter a struct we create a new file for the class defined by the struct. The class contains only public member fields.

The following struct

```
type point3d struct{
  x,y,z int;
}
```

gets translated to the following class file

```
.class public test1_struct_point3d
.super java/lang/Object

.field x I
.field y I
.field z I

.method public <init>()V
    .limit locals 99
    .limit stack 99
    aload_0
    invokenonvirtual java/lang/Object/<init>()V
    return
.end method
```

Append slices

Golite allows slices to be dynamically changing arrays using the `append()` function. Since jasmin does not have dynamically changing arrays we decided to have an algorithm which creates a new array with more space and then copies the values from the old array and finally, appends the new value. The golite code snippet corresponding to this is

```
var x [] int;
var newValue int= 100;
x = append(x, newValue);
```

The algorithm described in java is

```
int x[]; //old array
int y[]; //new array
int newValue = 100;
```

```
y = new int[x.length];
System.arraycopy(x,0,y,0,x.length);
y[x.length] = newValue;
x=y;
```

The target jasmin code is as described below. The actual algorithm is generic for any local and type.

```
aload_1 //old array ref in local 1
arraylength
iconst_1
iadd
newarray int
dup //duplicate new array reference
```

```

dup
dup
aload_1 // old array ref
swap
iconst_0
swap
iconst_0
aload_1
arraylength
invokestatic java/lang/System.arraycopy:(Ljava/lang/Object;ILjava/lang/Object;II)void
aload_1
arraylength
iload_2 //load new value
iastore //store in new array
astore_1 //store new array ref in local 1

```

If Else

If else statements are tricky in jasmin. Our approach is to evaluate the entire conditional expression and put it on the stack. Then we check if that expression is 1 (true) or false.

This code gen rule translates the below golite code

```

x := true
y := false

if y {
    //Do nothing
} else if x {
    //Do nothing
} else {
    println(x)
}

```

Into the following jasmin code

```

iconst_1
istore 0
iconst_0
istore 1
iload 1
ifne start1
goto stop0
start1:
iload_0
ifne start2
goto stop0

```

```

start2:
  getstatic  java/lang/System/out  Ljava/io/PrintStream;
  iload  0
  invokevirtual  java/io/PrintStream/println(Z)V
stop0:

```

Switch Statements

Jasmin has a native command for switch statements called lookup switch. However, golite allows switch cases to have multiple case expressions. For this we implemented an evaluation algorithm that turns switch statements into big if/else cases. The algorithm puts the switch case expression on the stack and has each case expression check for equality. If the check passes the code branches into the statement list otherwise it goes to the next case.

This code gen rule translates the below golite code

```

x := true
y := 1

switch y {
case 1,2:
  x = false;
default:
  x= true;
}

```

into the following jasmin code

```

  iconst_1
  istore 0
  ldc 1
  istore 1
  iload 1
  dup
  ldc 1
  ifeq startcase1
  dup
  ldc 2
  ifeq startcase1
goto stopcase1
startcase1:
  iconst_0
  istore 0
  goto endSwitch
stopcase1:
  iconst_1
  istore 0

```

```
endSwitch:
    pop
```

Loops

Loops are also similar to if/else statements except that we generate labels to the beginning and end of the loop. We pass the labels to the inner statement list such that a continue or break may goto the correct label. For any simple statements we evaluate them before the loop. Any increment statements get evaluated every loop at the end of the loop.

This code gen rule translates the below golite code

```
var y int;
for i:=0;i<10;i++){
    y=i+1
}
```

into the following jasmin code

```
        iconst_0
        istore 0
        ldc 0
        istore 1
start0:
        iload 1
        ldc 10
        if_icmplt true1
        iconst_0
        goto stop1
        true1:
        iconst_1
        stop1:
        ifne stop0

        iload 1
        ldc 1
        iadd
        istore 0

        iload 1
        iconst_1
        iadd
        istore 1
        goto start0
stop0:
```

For full code examples of source compiled to target see the TESTPROGRAMS directory.

3 Team Contribution

3.1 Shabbir Hussain

Shabbir contributed to the:

- Parser for the Statements, If/else, switch case rules.
- Milestone 1 report, Milestone 2 report, this document
- Milestone 1,2 test cases and testing
- Code Generation loops, statements, some expressions, append

3.2 Ossama Ahmed

For this project Sam served as the OCaml expert.

3.3 Michael Ho

4 Resources Used

Some of the resources we used are:

- <https://gobyexample.com>
- <http://caml.inria.fr/pub/docs/manual-ocaml-4.00/manual026.html>
- <https://realworldocaml.org/v1/en/html/parsing-with-ocamllex-and-menhir.html>

5 Conclusion and Future Work

Building a compiler is a very large undertaking and in this section we want to discuss some of the lessons learned.

One of the most important lessons learned is the value of automated testing. Our first deliverable was lacking due to the fact that we missed many test cases. On the second deliverable we spent time designing a small test framework for our Ocaml project. The framework would compile the project and then run the compiler on hundreds of code examples testing for syntax, typechecking and code generation. This came in very handy not only to verify our work but also to have regression testing as we added features to the compiler. One drawback to this system is that the output of our test is a text file which needs to be manually verified. The manual verification takes a lot of time and understanding of how the test scripts are run. In the future it would be advantageous to also have automated verification in the testing framework such that the output of running passing tests is a green light and failing tests point to which tests failed. Another lesson learned is having all modules depend on the same AST is a terrible design choice. One change in the AST requires a change in all the

dependant modules. In the future, it would be useful to design a tree search module which serves as an interface to the AST. This way, a change to the AST would only result in a change to the tree search module similar to our generator module. The reason we didn't implement this is because it was very easy to copy the tree searching code from the weeder to the pretty printer to the code-generator.